

# Machine Learning

CSE 142


Xin (Eric) Wang

Friday, October 8, 2021

**T  
o  
d  
a  
y**

- Classification, Ch. 2 & 3



- You, as a classifier designer, can often move decision boundaries (modify thresholds) to make the **false positive rate** as high or as low as you wish
  - A very high threshold (don't let anything through!) results in no **false positives** – but lots of **false negatives**
  - A very low threshold (let everything through!) results in no **false negatives** – but lots of **false positives**





- This doesn't necessarily make the classifier better or worse – it just changes the **operating point** of the classifier
- This is often application-specific:
  - When might false positives be especially undesirable?
  - When might false negatives be especially undesirable?
  - We can encode these preferences in a **cost function** to compute an optimal threshold, given this information

# Classifier design

A or B?



A or B?



It also depends on how good your **features** are!


- Either *raw* features or *constructed* features

# Feature separation vs. classifier design

---

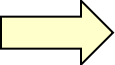
Placing the separating boundary = **classifier design**

Increasing the feature separation = **feature construction**



We can design a better classifier starting with the **features** in **C**!

# Typical predictive machine learning scenarios

<i>Task</i>	<i>Label space</i>	<i>Output space</i>	<i>Learning problem</i>
Classification	$\mathcal{L} = \mathcal{C}$	$\mathcal{Y} = \mathcal{C}$	learn an approximation $\hat{c} : \mathcal{X} \rightarrow \mathcal{C}$ to the true labelling function $c$
Scoring and ranking	$\mathcal{L} = \mathcal{C}$	$\mathcal{Y} = \mathbb{R}^{ \mathcal{C} }$	learn a model that outputs a score vector over classes
 Probability estimation	$\mathcal{L} = \mathcal{C}$	$\mathcal{Y} = [0, 1]^{ \mathcal{C} }$	learn a model that outputs a probability vector over classes
Regression	$\mathcal{L} = \mathbb{R}$	$\mathcal{Y} = \mathbb{R}$	learn an approximation $\hat{f} : \mathcal{X} \rightarrow \mathbb{R}$ to the true labelling function $f$

# Class probability estimation

---

A **class probability estimator** is a scoring classifier that outputs **probabilities** over the  $k$  classes – i.e., a mapping:

$$\hat{p} : \mathcal{X} \rightarrow [0, 1]^k$$

where

$$\sum_{i=1}^k \hat{p}_i(x) = 1$$

A key issue here is that we generally do not have access to the **true probabilities** for training data.

- E.g., an email is either spam or ham – it doesn't have a probability of being spam!
- So how can we train to learn such probabilities?

# Empirical probabilities

---

- In machine learning, we often calculate *empirical probabilities*
  - i.e., calculate **relative frequencies** from the available data

$N_i$  instances of the class  $C_i$  in the training data  $S$ :

$$\text{Relative frequency} = \frac{N_i}{|S|} = \hat{p}_i$$

- But this can be problematic, especially with small amounts of training data
  - Probabilities of 0 and 1 generally should be avoided
- There are various common ways to smooth or correct the relative frequencies to avoid 0 and 1
  - E.g., **Laplace correction** and **m-estimate**:

Add a pseudo-count to each class

$$\text{Laplace correction} = \frac{N_i + 1}{|S| + k}$$

Choose number of pseudo-counts  $m$  and their class distribution  $\pi_i$

$$\text{m-estimate} = \frac{N_i + m\pi_i}{|S| + m} \quad \sum_i \pi_i = 1$$

Think of  $\pi_i$  as the prior probabilities, and  $m$  the amount to weigh the priors relative to  $|S|$

# Quiz: empirical probabilities

---

Training data set  $S$

$C_1$ : 7 instances

$C_2$ : 14

$C_3$ : 0

$C_4$ : 4

$$\text{Relative frequency} = \frac{N_i}{|S|} = \hat{p}_i$$

$\hat{p}_1$ :

$\hat{p}_2$ :

$\hat{p}_3$ :

$\hat{p}_4$ :

$$\text{m-estimate} = \frac{N_i + m\pi_i}{|S| + m} \quad \sum_i \pi_i = 1$$

$$\text{Laplace correction} = \frac{N_i + 1}{|S| + k}$$

$m = 40 : \{10, 10, 10, 10\}$        $m = 20 : \{5, 0, 9, 6\}$



# Quiz: empirical probabilities

Training data set  $S$

$C_1$ : 7 instances

$C_2$ : 14

$C_3$ : 0

$C_4$ : 4

$$\text{Relative frequency} = \frac{N_i}{|S|} = \hat{p}_i$$

$$\begin{aligned} |S| &= 25 & \hat{p}_1 &: 7/25 = 0.28 \\ & & \hat{p}_2 &: 14/25 = 0.56 \\ & & \hat{p}_3 &: 0/25 = 0.0 \\ & & \hat{p}_4 &: 4/25 = 0.16 \end{aligned}$$

$$\text{m-estimate} = \frac{N_i + m\pi_i}{|S| + m} \quad \sum_i \pi_i = 1$$

$$\pi_i = \{0.25, 0.25, 0.25, 0.25\}$$

$$m = 40 : \{10, 10, 10, 10\}$$

$$\pi_i = \{0.25, 0, 0.45, 0.3\}$$

$$m = 20 : \{5, 0, 9, 6\}$$

$$\hat{p}_1: 17/65 = 0.26$$

$$\hat{p}_2: 24/65 = 0.37$$

$$\hat{p}_3: 10/65 = 0.15$$

$$\hat{p}_4: 14/65 = 0.22$$

$$\hat{p}_1: 12/45 = 0.27$$

$$\hat{p}_2: 14/45 = 0.31$$

$$\hat{p}_3: 9/45 = 0.2$$

$$\hat{p}_4: 10/45 = 0.22$$

$$\text{Laplace correction} = \frac{N_i + 1}{|S| + k}$$

$$\hat{p}_1: 8/29 = 0.28$$

$$\hat{p}_2: 15/29 = 0.52$$

$$\hat{p}_3: 1/29 = 0.03$$

$$\hat{p}_4: 5/29 = 0.17$$

# Multi-class classification – beyond binary!

---

- Many classification problems involve **multiple classes**
- Performance can be described with the **multi-class contingency table**
  - Also known as the **confusion matrix**
  - We can compute **accuracy**, per-class **precision**, per-class **recall**...

		<i>Predicted</i>			
<i>Actual</i>	15	2	3	20	
	7	15	8	30	
	2	3	45	50	
	24	20	56	100	

$$\text{Accuracy} = (15+15+45)/100 = 0.75$$


$$\text{Class 1 precision} = 15/24 = 0.63$$

$$\text{Class 1 recall} = 15/20 = 0.75$$


Etc.

# K-class classifiers

- How to build a *k*-class classifier?
  - We can combine several binary classifiers, e.g.:
    - **One-versus-rest scheme #1** – learn  $k-1$  models, apply in sequence
      - C1 vs. { C2, C3, C4 }
      - C2 vs. { C3, C4 }
      - C3 vs. C4




- **One-versus-rest scheme #2** – learn a *one-class* model for each class
  - C1 vs. { C2, C3, C4 }
  - C2 vs. { C1, C3, C4 }
  - C3 vs. { C1, C2, C4 }
  - C4 vs. { C1, C2, C3 }




# K-class classifiers

- **One-versus-one scheme #1** – learn a model for each pair of classes
  - Train  $k(k-1)/2$  binary classifiers, apply them all to  $x$  and **vote**




- **One-versus-one scheme #2** with a decision tree:



# Example: A 3-class linear classifier

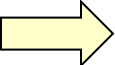
---

Classify instances into three classes  $\{A, B, C\}$  using three **linear discriminant functions** classifying (A vs. B), (A vs. C), and (B vs. C)



This implements the **one-versus-one scheme #2** on the previous slide

# Typical predictive machine learning scenarios

<i>Task</i>	<i>Label space</i>	<i>Output space</i>	<i>Learning problem</i>
Classification	$\mathcal{L} = \mathcal{C}$	$\mathcal{Y} = \mathcal{C}$	learn an approximation $\hat{c} : \mathcal{X} \rightarrow \mathcal{C}$ to the true labelling function $c$
Scoring and ranking	$\mathcal{L} = \mathcal{C}$	$\mathcal{Y} = \mathbb{R}^{ \mathcal{C} }$	learn a model that outputs a score vector over classes
Probability estimation	$\mathcal{L} = \mathcal{C}$	$\mathcal{Y} = [0, 1]^{ \mathcal{C} }$	learn a model that outputs a probability vector over classes
 Regression	$\mathcal{L} = \mathbb{R}$	$\mathcal{Y} = \mathbb{R}$	learn an approximation $\hat{f} : \mathcal{X} \rightarrow \mathbb{R}$ to the true labelling function $f$

# Regression – another predictive ML task


---

- In the classification tasks we've been discussing, the **label space** was a discrete set of classes
  - Classification, scoring, ranking, probability estimation
- **Regression** learns a function (the **regressor**) that is a mapping  $\hat{f} : \mathcal{X} \rightarrow \mathbb{R}$  from examples –  $f(x_i)$ 
  - I.e., the **target variable** (output) is real-valued
- Assumption: the examples will be noisy, so watch out for **overfitting** – we want to capture the general trend or shape of the function, not exactly match every data point
  - E.g., if fitting an **N-degree polynomial** to the training data (thus N+1 parameters to estimate), choose as small N as possible
- The number of data points should be much greater than the number of parameters to be estimated!
  - How much data is needed? This is an open question in ML....

# Regression example

Training data

$x$	$f(x)$
1.0	1.2
2.5	2.0
4.1	3.7
6.1	4.6
7.9	7.0



$\{1, 2, 3, 4, 5\}$ -degree polynomial functions

The regression function **may or may not** fit the training data exactly



# Regression

---

- We'll generally estimate a regression function based on some function of the *residual*, the different between the *estimate* and the *label* (the true value):

$$r(x) = f(x) - \hat{f}(x)$$

↑  
True function

↑  
Regression function

- That function is (again) the *loss function*  $L$ 
  - The most common loss function for regression is the *squared residual*:

$$L(r) = \mathbb{E} [r^2(x)] = \mathbb{E} [(f(x) - \hat{f}(x))^2]$$

# Bias-Variance Tradeoff

---

$$\begin{aligned}L(r) &= \mathbb{E} \left[ (f(x) - \hat{f}(x))^2 \right] \\&= \mathbb{E} \left[ (f(x) - \mathbb{E}[\hat{f}(x)] + \mathbb{E}[\hat{f}(x)] - \hat{f}(x))^2 \right] \quad (a+b)^2 = a^2 + b^2 + 2ab \\&= \mathbb{E} \left[ (f(x) - \mathbb{E}[\hat{f}(x)])^2 \right] + \mathbb{E} \left[ (\mathbb{E}[\hat{f}(x)] - \hat{f}(x))^2 \right] \\&\quad + \mathbb{E} \left[ 2(f(x) - \mathbb{E}[\hat{f}(x)]) (\mathbb{E}[\hat{f}(x)] - \hat{f}(x)) \right]\end{aligned}$$

$f(x)$  and  $\mathbb{E}[\hat{f}(x)]$  constant with respect to the noise

$$= (f(x) - \mathbb{E}[\hat{f}(x)])^2 + \mathbb{E} \left[ (\mathbb{E}[\hat{f}(x)] - \hat{f}(x))^2 \right]$$

$$= \text{Bias}^2 + \text{Variance}$$



*difference between the  
average prediction of our  
model and the true value*



*variations of our training  
data*


# Bias-Variance Tradeoff

$$L(r) = (f(x) - \mathbb{E}[\hat{f}(x)])^2 + \mathbb{E}[(\mathbb{E}[\hat{f}(x)] - \hat{f}(x))^2]$$


= Bias<sup>2</sup> + Variance

*difference between the  
average prediction of our  
model and the true value*


*variations of our training  
data*




Low bias, low variance



Low bias, high variance



High bias, low variance




High bias, high variance

# Bias-Variance Tradeoff


- Model with high variance pays a lot of attention to training data and does not generalize on the data which it hasn't seen before –

## Overfitting


- Such models perform very well on training data but poorly on test data
- Model with high bias pays very little attention to the training data and oversimplifies the model – **Underfitting**
  - High error on training and test data



overfitting



underfitting



Good balance

NEXT

## Concept Learning

READ

Chapter 4 in the textbook

*Logical Models: tree models and rule models*