

# Machine Learning

CSE 142

Xin (Eric) Wang

Friday, November 5, 2021

**T  
o  
d  
a  
y**


- SVM

# Notes

---

- HW2 due tonight
  - You should use your ucsc email and ucsc id as the username to register a CodaLab account otherwise we cannot know if you submit the code or not
  - Specify your username in your HW submission to Gradescope
  - You can choose your own teamname to show on the leaderboard for anonymity
- Midterm grades will be out next week

- The **margin ( $z$ )** of a sample is its **distance** from the classification boundary
  - Positive if it's correctly classified
  - Negative if it's incorrectly classified




Perceptron margin for point  $\mathbf{x}$ :

$$\mathbf{z}(\mathbf{x}) = \frac{y(\mathbf{w}^T \mathbf{x} - t)}{\|\mathbf{w}\|} = \frac{m}{\|\mathbf{w}\|} \quad \text{Non-homogeneous representation}$$

Note:  $m$  is not the **margin**; it's the result of plugging  $\mathbf{x}_i$  into  $y(\mathbf{w}^T \mathbf{x} - t)$

- The **class margin** (on the training set) is the **minimum margin** of the data points for that class
- The **classifier margin** is the sum of the class margins
- There are an infinite number of linear classifiers that can perfectly separate linearly separable data




Infinitely many lines can separate the two classes

- But which (of all these) is the **best** linear classifier?
  - Perhaps the one that **maximizes the classifier margin**

# Support vector machine (SVM)

---

- A **support vector machine (SVM)** is a linear classifier whose decision boundary is a linear combination of the **support vectors** (training samples at the margins)
- In an SVM, we find classifier parameters  $(\mathbf{w}, t)$  that **maximize the classifier margin**
- Since  $m = y(\mathbf{w}^T \mathbf{x} - t)$  and we wish to **maximize the margin**  $\frac{m}{\|\mathbf{w}\|}$ , we can instead fix  $m = 1$  and **minimize  $\|\mathbf{w}\|$** 
  - Provided that **none of the training points fall inside the margin**
- This leads to a **constrained optimization problem**:

$$\mathbf{w}^*, t^* = \operatorname{argmin}_{\mathbf{w}, t} \frac{1}{2} \|\mathbf{w}\|^2 \quad \text{subject to } y_i(\mathbf{w} \cdot \mathbf{x}_i - t) \geq 1, 1 \leq i \leq n$$


- Then, after some quadratic optimization based on Lagrange multiplier (Page 212-214)....

# Support vector machine (SVM)

---

...we get the following result:

$$\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i \quad \text{where } \alpha_i \text{ are non-negative reals s.t. } \sum_{i=1}^n \alpha_i y_i = 0$$

$\alpha_i > 0$  only for the **support vectors**!

Other data  $\mathbf{x}_i$  for which  $\alpha_i = 0$  can be **removed from the training set** without affecting the learned decision boundary

I.e., the decision boundary is defined only by the (typically few) **support vectors** from the training set – those that are nearest to the decision boundary (at the margin)

And thus the weight vector  $\mathbf{w}$  is merely a linear combination of the (typically few) **support vectors**

The threshold  $t$  can be found by solving  $m = 1 = \mathbf{w}^T \mathbf{x} - t$  for any **support vector**  $\mathbf{x}$

# Support vector machine (SVM) in dual form

- How do we find the  $\alpha_i$  values?
  - Dual form of the optimization—a function of Lagrange multipliers only
  - Via a quadratic optimization solver!
  - In some simple problems we can do them by hand

$$\alpha_1^*, \dots, \alpha_n^* = \underset{\alpha_1, \dots, \alpha_n}{\operatorname{argmax}} \left[ -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i \cdot \mathbf{x}_j + \sum_{i=1}^n \alpha_i \right]$$

Note the pairwise dot products between training instances—the entries of the Gram matrix

$$\text{subject to } \alpha_i \geq 0, 1 \leq i \leq n \text{ and } \sum_{i=1}^n \alpha_i y_i = 0$$

1. Quadratic optimization to solve for  $\alpha_1, \dots, \alpha_n$ 
  - Non-zero  $\alpha_i$  corresponds to support vector  $\mathbf{x}_i$
2. Create  $\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$
3. Solve for  $t$  by plugging in for any support vector  $\mathbf{x}_i$ 
$$m = 1 = \mathbf{w}^T \mathbf{x}_i - t$$

The support vectors  $\mathbf{x}_i$  fully determine the decision boundary!

# SVM classifier example (Fig. 7.8)

$$\mathbf{X} = \begin{pmatrix} 1 & 2 \\ -1 & 2 \\ -1 & -2 \end{pmatrix} \quad \mathbf{y} = \begin{pmatrix} -1 \\ -1 \\ +1 \end{pmatrix} \quad \mathbf{X}' = \begin{pmatrix} -1 & -2 \\ 1 & -2 \\ -1 & -2 \end{pmatrix}$$


- The matrix  $\mathbf{X}'$  incorporates the class labels: the rows are  $y_i \mathbf{X}_i$
- The Gram matrix is:

$$\mathbf{X}\mathbf{X}^T = \begin{pmatrix} 5 & 3 & -5 \\ 3 & 5 & -3 \\ -5 & -3 & 5 \end{pmatrix} \quad \mathbf{X}'\mathbf{X}'^T = \begin{pmatrix} 5 & 3 & 5 \\ 3 & 5 & 3 \\ 5 & 3 & 5 \end{pmatrix}$$

- The dual optimization problem is thus

$$\begin{aligned} & \arg\max_{\alpha_1, \alpha_2, \alpha_3} -\frac{1}{2} \left( 5\alpha_1^2 + 3\alpha_1\alpha_2 + 5\alpha_1\alpha_3 + 3\alpha_2\alpha_1 + 5\alpha_2^2 + 3\alpha_2\alpha_3 + 5\alpha_3\alpha_1 + 3\alpha_3\alpha_2 + 5\alpha_3^2 \right) + \alpha_1 + \alpha_2 + \alpha_3 \\ & = \arg\max_{\alpha_1, \alpha_2, \alpha_3} -\frac{1}{2} \left( 5\alpha_1^2 + 6\alpha_1\alpha_2 + 10\alpha_1\alpha_3 + 5\alpha_2^2 + 6\alpha_2\alpha_3 + 5\alpha_3^2 \right) + \alpha_1 + \alpha_2 + \alpha_3 \end{aligned}$$

subject to  $\alpha_1 \geq 0, \alpha_2 \geq 0, \alpha_3 \geq 0$  and  $-\alpha_1 - \alpha_2 + \alpha_3 = 0$ .





# SVM classifier example (Fig. 7.8)

---

✚ Using the equality constraint we can eliminate one of the variables, say  $\alpha_3$ , and simplify the objective function to

$$\arg \max_{\alpha_1, \alpha_2, \alpha_3} -\frac{1}{2} (20\alpha_1^2 + 32\alpha_1\alpha_2 + 16\alpha_2^2) + 2\alpha_1 + 2\alpha_2$$

✚ Setting partial derivatives to 0 we obtain  $-20\alpha_1 - 16\alpha_2 + 2 = 0$  and  $-16\alpha_1 - 16\alpha_2 + 2 = 0$  (notice that, because the objective function is quadratic, these equations are guaranteed to be linear).


✚ We therefore obtain the solution  $\alpha_1 = 0$  and  $\alpha_2 = \alpha_3 = 1/8$ . We then have  $\mathbf{w} = 1/8(\mathbf{x}_3 - \mathbf{x}_2) = \begin{pmatrix} 0 \\ -1/2 \end{pmatrix}$ , resulting in a margin of  $1/||\mathbf{w}|| = 2$ .

✚ Finally,  $t$  can be obtained from any support vector, say  $\mathbf{x}_2$ , since  $y_2(\mathbf{w} \cdot \mathbf{x}_2 - t) = 1$ ; this gives  $-1 \cdot (-1 - t) = 1$ , hence  $t = 0$ .

# SVM classifier example (Fig. 7.8)

We now add an additional positive at (3, 1):

$$\mathbf{X}' = \begin{pmatrix} -1 & -2 \\ 1 & -2 \\ -1 & -2 \\ 3 & 1 \end{pmatrix} \quad \mathbf{X}'\mathbf{X}'^T = \begin{pmatrix} 5 & 3 & 5 & -5 \\ 3 & 5 & 3 & 1 \\ 5 & 3 & 5 & -5 \\ -5 & 1 & -5 & 10 \end{pmatrix}$$



👉 It can be verified by similar calculations to those above that the margin decreases to 1 and the decision boundary rotates to  $\mathbf{w} = \begin{pmatrix} 3/5 \\ -4/5 \end{pmatrix}$ .


👉 The Lagrange multipliers now are  $\alpha_1 = 1/2$ ,  $\alpha_2 = 0$ ,  $\alpha_3 = 1/10$  and  $\alpha_4 = 2/5$ . Thus, only  $\mathbf{x}_3$  is a support vector in both the original and the extended data set.

# Support vector machine (SVM)

---

What if the data is not linearly separable?

Or a data point “strays” into an otherwise nice margin?




This can be solved with  
a **Soft Margin SVM**

# Soft margin SVM

$$\xi_i = \text{"xi"}$$

- We introduce a **slack variable**  $\xi_i$  for each training example to account for **margin errors**
  - Points that are **inside** the margin
  - Points that are on the **wrong side** of the decision boundary

$$\mathbf{w}^T \mathbf{x}_i - t \geq 1 - \xi_i$$




- Results in the soft margin optimization problem:

$$\mathbf{w}^*, t^*, \xi_i^* = \underset{\mathbf{w}, t, \xi_i}{\operatorname{argmin}} \left[ \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i \right]$$


subject to  $y_i(\mathbf{w} \cdot \mathbf{x}_i - t) \geq 1 - \xi_i$  and  $\xi_i \geq 0, 1 \leq i \leq n$

- The **complexity parameter**  $C$  is a **user-defined parameter** that allows for a tradeoff between **maximizing the margin** (lower  $C$ ) and **minimizing the margin errors** (higher  $C$ )
  - A high value of  $C$  means that margin errors incur a high penalty
  - A low value permits more margin errors (possibly including misclassifications) in order to achieve a large margin
  - Note that when  $C = 0$ , this gives no penalty to outliers – which makes it equivalent to our **basic linear classifier**!*

# Soft margin SVM



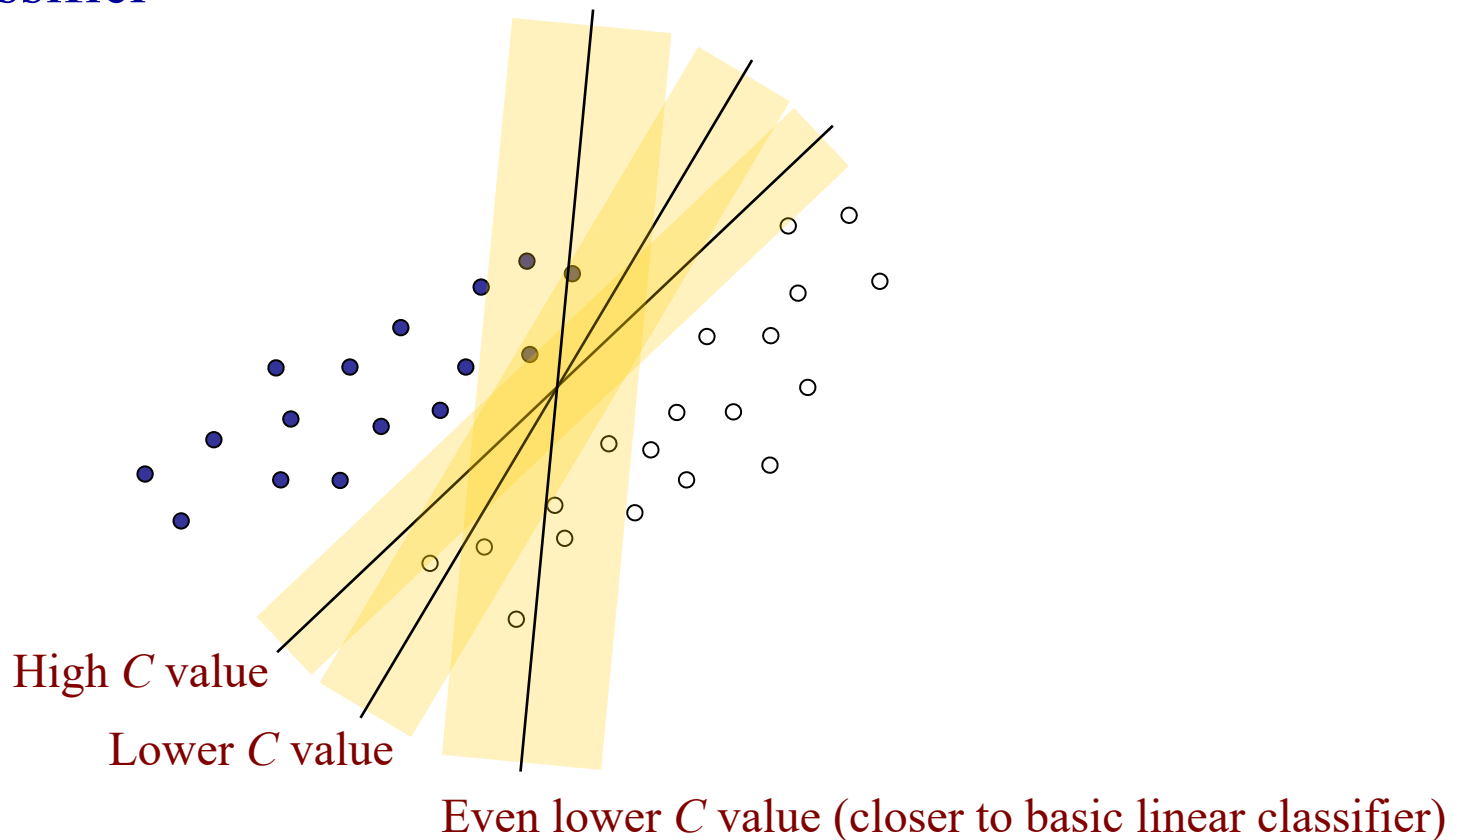
$$C = \frac{5}{16} \quad \begin{array}{l} \text{Smaller margin} \\ \text{Fewer margin errors} \end{array}$$



$$C = \frac{1}{10} \quad \begin{array}{l} \text{Larger margin} \\ \text{More margin errors} \end{array}$$

# Soft margin SVM

A minimal-complexity (low  $C$ ) soft margin classifier summarizes the classes by their class means in a way very similar to [the basic linear classifier](#)



# Quiz: Perceptron and SVM

---

- For both perceptron and SVM classifiers,  $\mathbf{w} = \sum_i \alpha_i y_i \mathbf{x}_i$ , how are  $\alpha_i$  and  $\mathbf{x}_i$  different in perceptron and SVM?



# Perceptron and SVM binary classifiers – summary

- In the **perceptron** model, we iteratively learn the linear discriminant  $\mathbf{w}$ , which is a linear combination of the misclassified input vectors  $\mathbf{x}_i$

$$\mathbf{w} = \sum_i \alpha_i y_i \mathbf{x}_i$$

$\alpha_i$  – # of times  $\mathbf{x}_i$  was misclassified  
 $y_i$  – class label of  $\mathbf{x}_i$   $\{+1, -1\}$

- After training, a new input is classified as a member of the positive class if  $\mathbf{w}^T \mathbf{x} > 0$  (using homogeneous representation)
- In **SVM** learning, we solve a constrained optimization problem:

$$\alpha_1^*, \dots, \alpha_n^* = \underset{\alpha_1, \dots, \alpha_n}{\operatorname{argmax}} \left[ -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i \cdot \mathbf{x}_j + \sum_{i=1}^n \alpha_i \right]$$

subject to  $\alpha_i \geq 0, 1 \leq i \leq n$  and  $\sum_{i=1}^n \alpha_i y_i = 0$

which leads us to  $\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$  where  $\alpha_i = 0$  except for the support vectors

Non-homogeneous

# Perceptron and SVM binary classifiers – summary

---

- In both perceptron and SVM learning, the linear decision boundary is a **linear combination of the training data points**
  - In the perceptron, just the ones that get **misclassified** in the iterative training
  - In the SVM, just the (few) **support vectors**
- Both learning methods have a **dual form** in which the **dot product** of training data points  $\mathbf{x}_i^T \mathbf{x}_j$  is part of the main computation
  - All values of  $\mathbf{x}_i^T \mathbf{x}_j$  are contained in the **Gram matrix**

$$\mathbf{G} = \mathbf{X}^T \mathbf{X} = [\mathbf{x}_1 \quad \mathbf{x}_2 \quad \dots \quad \mathbf{x}_k]^T [\mathbf{x}_1 \quad \mathbf{x}_2 \quad \dots \quad \mathbf{x}_k]$$

so it's often efficient to **compute the Gram matrix in advance** and index into it, rather than computing the dot products over and over again

# Perceptron and SVM binary classifiers – summary

---

- Perceptron and (basic) SVM learning only converge to a solution if the training data is **linearly separable**
- If the data is not linearly separable, we can employ a **soft margin SVM**, where we introduce a *slack variable*  $\xi_i$  for each training data point, allowing for margin errors:

$$\mathbf{w}^T \mathbf{x}_i - t \geq 1 - \xi_i \quad \xi_i > 0 \rightarrow \mathbf{x}_i \text{ is not a support vector}$$

and leading to this optimization problem:

$$\mathbf{w}^*, t^*, \xi_i^* = \underset{\mathbf{w}, t, \xi_i}{\operatorname{argmin}} \left[ \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i \right]$$

subject to  $y_i(\mathbf{w} \cdot \mathbf{x}_i - t) \geq 1 - \xi_i$  and  $\xi_i \geq 0, 1 \leq i \leq n$

where the **complexity parameter**  $C$  is a user-defined parameter that allows for a tradeoff between maximizing the margin (lower  $C$ ) and minimizing the margin errors (higher  $C$ )