

# Machine Learning

CSE 142

Xin (Eric) Wang

Monday, November 22, 2021

**T  
o  
d  
a  
y**

- Model ensembles (cont.) (Ch. 11)
- ML experiments (Ch. 12)

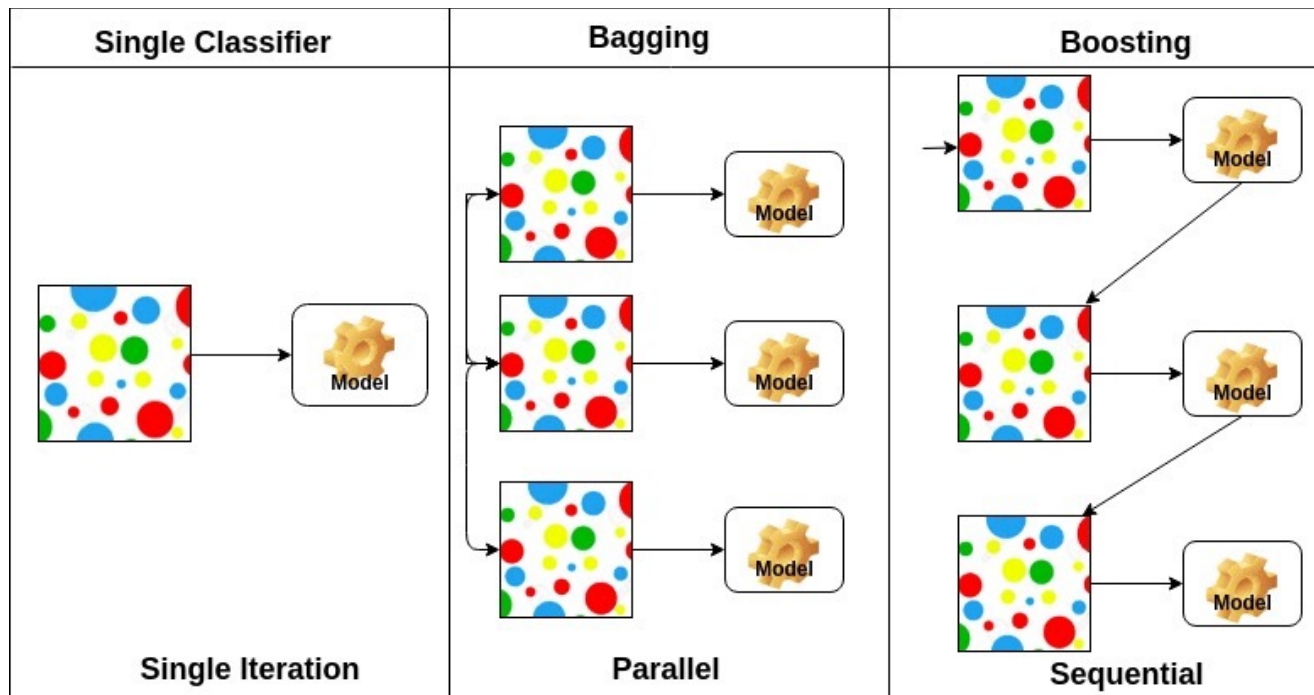
# Model ensembles

Chapter 11 in the textbook


- We've seen how **combining features** can be beneficial
- We can also **combine models** to increase performance
  - Combinations of models are known as model ensembles
  - Potential for better performance at the cost of increased complexity
- General approach to **model ensembles**:
  - **Construct** multiple different models from adapted versions of the training data (e.g., reweighted or resampled)
  - **Combine** the predictions of these models in some way (averaging, voting, weighted voting, etc.)
- Two of the best-known ensemble methods are **bagging** and **boosting**
- These are “**meta**” **methods** – i.e., they are independent of the particular learning method (linear classifier, SVM, etc.)

- Bagging = “bootstrap aggregation”
  - Create  $T$  models on different random samples of the training data set
  - Each sample is a set of training data called a bootstrap sample
    - Could be any size – often  $|D|$  is used, the size of the training set
- Bootstrap samples: Sample the data set with replacement (i.e., a data point can be chosen more than once in a bootstrap sample)
  - For  $j = 1$  to  $|D|$ , choose with uniform probability from training data points  $D = \{ d_1, d_2, \dots, d_{|D|} \}$
  - Gather these into the bootstrap sample  $D_i$
- Use  $\{ D_1, D_2, \dots, D_T \}$  to train models  $\{ M_1, M_2, \dots, M_T \}$ , then combine the predictions of the  $T$  models
- Differences between the  $T$  bootstrap samples create diversity among the models in the ensemble

- **Boosting** is similar to bagging, but it uses a more sophisticated technique to **create diverse training sets**
- The focus is on adding classifiers that do better on the **misclassifications** from earlier classifiers
  - By giving them a **higher weight**



- How much should the weights of training examples change?
  - Assign half of the total weights to the misclassified examples
- How to combine the individual models?
  - Confidence factor  $\alpha_t$  for each model
  - High  $\alpha_t$  to model with low error rate



# Boosting

Assumption is complex problem, lots of data, no perfect classifier. I.e.,  $\epsilon_t > 0$

- Procedure:


- Assign equal weights  $w_j$  to training data points
- Train a classifier; assign it a **confidence factor**  $\alpha_t$  based on the error rate  $\epsilon_t$

$$\alpha_t = \frac{1}{2} \ln \frac{1-\epsilon_t}{\epsilon_t}$$

- Give **misclassified** instances a **higher weight**
  - Assign **half** of the total weight to the misclassified examples
- Repeat for  $T$  classifiers or until  $\epsilon_t \geq 0.5$
- The ensemble predictor is a **weighted average** of the models (rather than majority vote)

$$M(x) = \sum_{t=1}^T \alpha_t M_t(x)$$

- Threshold for binary output



Misclassified  
points

$$w' = \frac{w}{2\epsilon_t}$$

Correctly classified  
points

$$w' = \frac{w}{2(1 - \epsilon_t)}$$

Weights will sum to  $0.5+0.5=1$

# Boosting

Computational complexity:  $O(TDK)$   
(ensemble size x data size x dimensionality)

---

**Algorithm** Boosting( $D, T, \mathcal{A}$ ) – train an ensemble of binary classifiers from reweighted training sets.

---

**Input** : data set  $D$ ; ensemble size  $T$ ; learning algorithm  $\mathcal{A}$ .

**Output** : weighted ensemble of models.

```
 $w_{1i} \leftarrow 1/|D|$  for all  $x_i \in D$  ; // start with uniform weights
for  $t = 1$  to  $T$  do
    run  $\mathcal{A}$  on  $D$  with weights  $w_{ti}$  to produce a model  $M_t$ ;
    calculate weighted error  $\epsilon_t$ ;
    if  $\epsilon_t \geq 1/2$  then
        | set  $T \leftarrow t - 1$  and break
    end
     $\alpha_t \leftarrow \frac{1}{2} \ln \frac{1-\epsilon_t}{\epsilon_t}$  ; // confidence for this model
     $w_{(t+1)i} \leftarrow \frac{w_{ti}}{2\epsilon_t}$  for misclassified instances  $x_i \in D$  ; // increase weight
     $w_{(t+1)j} \leftarrow \frac{w_{tj}}{2(1-\epsilon_t)}$  for correctly classified instances  $x_j \in D$  ; // decrease
end
return  $M(x) = \sum_{t=1}^T \alpha_t M_t(x)$ 
```

---



# Boosting example

Misclassified points


$$w' = \frac{w}{2\epsilon_t}$$

Correctly classified points

$$w' = \frac{w}{2(1 - \epsilon_t)}$$

29 points

Initial weights  $w = 1/29 = 0.034$



Round 1:

6 errors

$$\epsilon_1 = \frac{6}{29} = 0.21$$

$$\alpha_1 = \frac{1}{2} \ln \frac{1 - \epsilon_1}{\epsilon_1} = 0.67$$

$$w' = \left\{ \frac{w}{2\epsilon_t}, \frac{w}{2(1 - \epsilon_t)} \right\} = \{ 0.082, 0.021 \}$$

Round 2:

7 errors  $\Rightarrow \epsilon_2 = \frac{7}{29} = 0.24$  ?

$$\epsilon_2 = \frac{2 * 0.082 + 5 * 0.021}{1} = 0.28$$


$$\alpha_2 = \frac{1}{2} \ln \frac{1 - \epsilon_2}{\epsilon_2} = 0.47$$

$$w' = \left\{ \frac{w}{2\epsilon_t}, \frac{w}{2(1 - \epsilon_t)} \right\}$$

Continue for  $T$  iterations or until  $\epsilon \geq 0.5$


# Boosting vs. bagging

Boosting



Can achieve zero training error by focusing on the misclassifications

Bagging



With relatively large bootstrap sample sets, there tends to be little diversity in the learned models

# Ensemble methods: summary

---

- Model ensemble methods are **meta-methods**: they are ways to combine ML models, but they are agnostic to the kind of model being used
  - Simple linear classifier, Perceptron, SVM, neural network, etc. (weighted versions)
- They provide an opportunity to **improve performance** and (mostly) **avoid over-fitting** to the training data
  - Typically using randomization of some sort
- Bagging is essentially a **variance reduction** technique (**increase consistency**), while boosting is essentially a **bias reduction** technique (**increase accuracy**)
- There are many other **ensemble** approaches in machine learning

# Machine learning experiments

## Chapter 12

# Machine learning experiments

---

- While there is a lot of important **theory** in machine learning (e.g., algorithms, proving convergence, proving performance bounds, establishing computability, etc.), it is a practical field, applying algorithms to data for practical use
- ML **experiments** are critical for confirming the various assumptions in a ML problem and establishing **realistic expectations for performance**
- Some key questions for ML experiments in a domain  $D$ :
  - How does a specific model perform on data from  $D$ ?
  - Which of a set of models has the best performance on  $D$ ?
  - How do models from learning algorithm  $A$  perform on  $D$ ?
  - Which learning algorithm  $A$  produces the best model(s) for domain  $D$ ?
- How do we define **performance**?

$$\text{False positive rate (FPR)} = \frac{FP}{N} = \alpha$$

$$\text{Accuracy} = \frac{TP+TN}{P+N} = \left(\frac{P}{P+N}\right)TPR + \left(\frac{N}{P+N}\right)TNR$$

$$\text{False negative rate (FNR)} = \frac{FN}{P} = \beta$$

$$\text{Error rate} = \frac{FP+FN}{P+N}$$

$$\text{True positive rate (TPR)} = \frac{TP}{P} = \text{Sensitivity} = \text{Recall} = 1 - \beta$$

$$\text{Precision} = \frac{TP}{\hat{P}}$$

$$\text{True negative rate (TNR)} = \frac{TN}{N} = \text{Specificity} = 1 - \alpha$$

$$\text{Accuracy} + \text{error rate} = 1$$

$$\text{Average recall} = \frac{TPR+TNR}{2}$$

Contingency table:

		Actual class $C$		
		1	0	
Predicted class $\hat{C}$	1	TP	FP	<i>Estimated positive <math>\hat{P}</math></i>
	0	FN	TN	<i>Estimated negative <math>\hat{N}</math></i>
		<i>Positives <math>P</math></i>	<i>Negatives <math>N</math></i>	TOTAL

# Example

---

		$C$	
		1	0
$\hat{C}$	1	50	0
	0	50	10
		P=100	N=10

$$\text{Accuracy} = 60/110 = 0.55$$

$$\text{Error rate} = 1 - 0.55 = 0.45$$

$$\text{Average recall} = (0.5+1.0)/2 = 0.75$$

		$C$	
		1	0
$\hat{C}$	1	50	0
	0	50	100
		P=100	N=100

$$\text{Accuracy} = 150/200 = 0.75$$

$$\text{Error rate} = 1 - 0.75 = 0.25$$

$$\text{Average recall} = (0.5+1.0)/2 = 0.75$$

# Performance measure

---

- If we choose **accuracy** (or **error rate**) as the measure of performance, we're making an assumption that the **class distribution** of the test set is **representative** of the real problem domain
- If we choose **average recall**, we're making an assumption that the real problem consists of **uniform class distributions** (equally likely)
- In general, when running an experiment, we should keep a record of a sufficient set of measurements to enable reproduction of the **contingency table** – in order to compute additional values later on if needed
  - Need four of the independent values (e.g., FP, FN, P, N)



# Performance measure (testing)

---

- We want to estimate the **true error rate** = the error rate on the entire population (not just on your data set)
- You cannot use your full dataset to **train** a model, estimate the model **parameters**, and estimate the **performance** (e.g., accuracy or error rate) of your model
  - This results in **overfitting** and overly **optimistic** results!
- **Testing data** is used to predict the performance of your ML model on the real problem
  - Applied **after the model is finalized**; it is not used to modify the model
  - In other words, **testing** is just a substitute for deploying your model in a product and gathering data from millions of users
  - Only **training** and **validation** data can change your model (\*)


# Training, validation, and testing

---

- Testing on the training data, or training on the testing data, is **cheating!**
  - **Training/validation** data and **testing** data must be completely separate
  - The **test data** is (conceptually) stored in a vault until ready to be used for testing the completed model
  - You **cannot modify the model** after assessing it with the testing data
    - At least, you can't modify it and re-test
- Instead, we split the training data into **disjoint** sets for validation
  - **Holdout method**
    - Keep some data separate (typically 10-30%) for testing
  - **k-fold cross-validation**: average the  $k$  different performance estimates
  - **Leave-one-out cross-validation**
    - When  $k =$  the number of data points

# Typical training/validation/testing example

- For a **1D regression problem**, let's train different models (different regression functions):
  - 0-degree polynomial (a constant), 1-degree polynomial (a line), 2<sup>nd</sup>-degree polynomial, 3<sup>rd</sup>-degree polynomial, 4<sup>th</sup>-degree polynomial
- Errors on the <sup>(blue dots)</sup>training data:



Which is the best model?

Answer: We don't know!

Use the <sup>(red dots)</sup>validation data to determine:  
The **second-degree polynomial** has the lowest error on this validation data

Now run **cross-validation** on the data set to estimate the performance of your model

Produce the **final model** by using all your data


Ship it!

# Cross-validation

---

- The **cross-validation** process for experimental evaluation:
  - Randomly partition the experimental data into  $k$  parts (“folds”)
  - **Train** the model on  $k-1$  folds
  - **Evaluate** it on the remaining test fold
  - **Repeat** for a total of  $k$  times, evaluating on each test fold
  - **Average** the performance measure (e.g., accuracy), over the  $k$  trials
- This is  **$k$ -fold cross-validation**
  - Frequently,  $k = 10$ , but other values are also used
    - $k = 2$  or  $5$  is quite common
  - Rule of thumb: folds should contain **at least 30 instances**
    - Implying  $\geq 300$  instances in the data set for  $k = 10$
  - If  $k = n$  (the number of data points), this is **leave-one-out cross-validation**

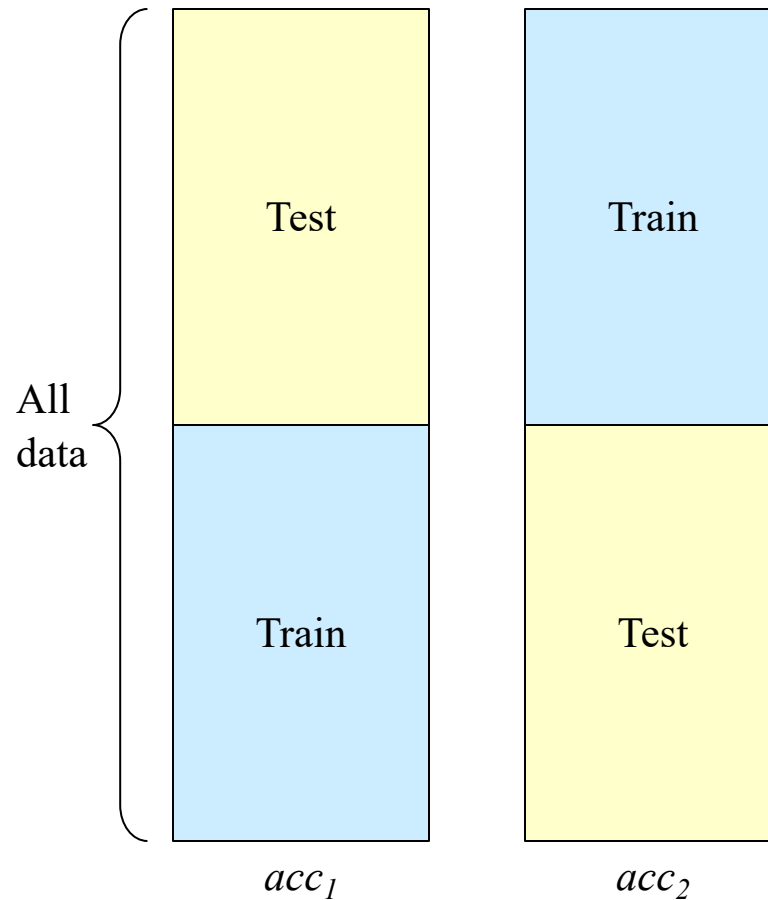
# Example: 10-fold cross-validation



$$acc = \text{Average}(acc_i)$$

# Example: 2-fold cross-validation

---




$$acc = \frac{1}{2} (acc_1 + acc_2)$$

# Repeated random sub-sampling validation

**Randomly** split the data into  $k$  folds  $N$  times, and then average the  $N$  results of  $k$ -fold cross-validation


- Typically,  $k = 2$



$acc_1$

$acc_2$


$$acc_a = \frac{1}{2} (acc_1 + acc_2)$$



$acc_1$

$acc_2$

$$acc_b = \frac{1}{2} (acc_1 + acc_2)$$



$acc_1$

$acc_2$

$$acc_c = \frac{1}{2} (acc_1 + acc_2)$$

$$\hat{a} = (acc_a + acc_b + acc_c)/3$$

# Cross-validation

---

- In practice, the choice of the number of **fold**s depends on the size of the data set
  - For **large** data sets, fewer folds are needed
  - For **sparse** data sets, **leave-one-out** may be best
- We should have some notion of **acceptable variance** for the cross-validation performance measure
  - If the variance is high (among the  $k$  performance measures), we probably need more data!
- If we are satisfied with the performance of our learning algorithm, we then run it over the entire data set (i.e., train on the complete data set) to **produce the final model (\*)**
  - Hence the term **cross-validation** rather than cross-testing!
  - (In this case – oddly – we never test the *final* model.)



# Notes on terminology

---

- Train (training), validate (validation), test (testing), evaluate (evaluation)
  - There can be some ambiguity among these term, depending on the context
- Keep in mind:
  - We can **evaluate** a candidate model
    - **Validation**
  - We can **evaluate** the final system performance
    - **Testing**
  - We can **evaluate** performance on the training data
    - This might be done as part of the **training** algorithm or as part of the **validation** process
  - A given data point (from your data set) can play multiple roles (training, validation, and testing)
    - But never use data for **validation or testing** that was used for **training** on the same model!

# Model parameters and hyperparameters

---

- Machine learning algorithms **learn the model parameters**
  - You choose what **model** is to be learned (linear regression, decision tree, SVM, neural network, etc.)
  - Training sets the values of various **parameters** (e.g., the  $\mathbf{w}$  vector, the DT decision nodes, the  $\alpha$  values for a kernel perceptron, etc.)
- Models may also have **hyperparameters** – settings or parameters that are determined outside the learning algorithm itself, fixed during training
  - E.g., the kernel function used, the highest order of a polynomial function, the value of  $k$  for  $k$ -NN, the number of hidden layers in a NN, etc.
- We can think of **validation** as the process of determining the preferred **hyperparameters**
  - Train to set **parameters**; validate to determine **hyperparameters**

# Quiz: ML experiments

---

- See the quiz “ML experiments” on Canvas